

```
In [1]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [2]: data_dir = '/content/drive/My Drive/sample-fruits-360'
```

```
In [3]: # Import Libraries
import numpy as np
import tensorflow as tf
from tensorflow import keras
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
from torch.autograd import Variable
from torch.utils.data import DataLoader
from sklearn.datasets import load_files
from sklearn.model_selection import train_test_split
```

```
In [4]: # Function for Load images
def load_dataset(path):
    data = load_files(path)
    files = np.array(data['filenames'])
    targets = np.array(data['target'])
    target_labels = np.array(data['target_names'])
    return files, targets, target_labels
```

```
In [5]: # Load Dataset
x, y, target_labels = load_dataset(data_dir)
print("Dataset Loaded !")

# Get Training size and Test size
print('Total set size : ',x.shape)
print('Total targets : ',len(target_labels) )
```


Dataset Loaded !
Total set size : (1095,)
Total targets : 73

```
In [6]: # Function for convert image to array
def convert_image_to_array(files):
    images_as_array=[]
    for file in files:
        images_as_array.append(keras.preprocessing.image.img_to_array(keras.pr
eprocessing.image.load_img(file)))
    return images_as_array

# Convert images to numpy array using keras.preprocessing library
x = np.array(convert_image_to_array(x),np.float32)
print(x.shape)

(1095, 100, 100, 3)
```

```
In [7]: # Plot image on random data
plt.imshow(x[1]/255)
plt.title(target_labels[y[1]])
plt.show()
```



```
In [8]: # Flatten the features of image
x = x.reshape([-1,100*100*3])
x = x/255
print("final shape : " , x.shape)
```

final shape : (1095, 30000)

```
In [9]: # Train and Test split
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=74)

# Get size of all set
print("X Train size : ", X_train.shape)
print("X Test size : ", X_test.shape)
print("Y Train size : ", y_train.shape)
print("Y Test size : ", y_test.shape)
```

X Train size : (876, 30000)

X Test size : (219, 30000)

Y Train size : (876,)

Y Test size : (219,)

```
In [10]: # Convert numpy array to torch
X_train = torch.from_numpy(X_train)
y_train = torch.from_numpy(y_train).type(torch.LongTensor)
X_test = torch.from_numpy(X_test)
y_test = torch.from_numpy(y_test).type(torch.LongTensor)
```

```
In [11]: # Define no of iteration, batch size, num_epochs
batch_size=100
n_iters = 1000
num_epochs = n_iters / (len(X_train) / batch_size)
num_epochs = int(num_epochs)
```

```
In [12]: # Set train and test
train = torch.utils.data.TensorDataset(X_train,y_train)
test = torch.utils.data.TensorDataset(X_test,y_test)
train_loader = DataLoader(train, batch_size = batch_size, shuffle = False)
test_loader = DataLoader(test, batch_size = batch_size, shuffle = False)
```

CNN with 2 convolutional layer and 1 fully connected layer

```
In [13]: # Create CNN Model
class CNNModel(nn.Module):
    def __init__(self):
        super(CNNModel, self).__init__()

        self.cnn1 = nn.Conv2d(in_channels=3,out_channels=16,kernel_size=5,stride=1,padding=0)
        self.relu1 = nn.ReLU()

        self.maxpool1 = nn.MaxPool2d(kernel_size=2)

        self.cnn2 = nn.Conv2d(in_channels=16,out_channels=32,kernel_size=5,stride=1,padding=0)
        self.relu2 = nn.ReLU()

        self.maxpool2 = nn.MaxPool2d(kernel_size=2)

        self.fc1 = nn.Linear(15488,len(target_labels));

    def forward(self,x):

        out=self.cnn1(x)
        out=self.relu1(out)
        out=self.maxpool1(out)

        out=self.cnn2(out)
        out=self.relu2(out)
        out=self.maxpool2(out)

        out = out.view(out.size(0), -1)
        out=self.fc1(out)

        return out
```

```
In [14]: # Initialize Parameters and fit the model
model = CNNModel()
error = nn.CrossEntropyLoss()
learning_rate = 0.1
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```

```

In [15]: # CNN model training
count = 0
loss_list = []
iteration_list = []
accuracy_list = []
for epoch in range(num_epochs):
    for i, (images, labels) in enumerate(train_loader):
        train = Variable(images.view(-1,3,100,100))
        #print(train.shape)
        labels = Variable(labels)
        optimizer.zero_grad()
        outputs = model(train)
        loss = error(outputs, labels)
        loss.backward()
        optimizer.step()
        count += 1

    if count % 50 == 0:
        correct = 0
        total = 0
        for images, labels in test_loader:
            test = Variable(images.view(-1,3,100,100))
            outputs = model(test)
            predicted = torch.max(outputs.data, 1)[1]
            total += len(labels)
            correct += (predicted == labels).sum()
        accuracy = 100 * correct / float(total)
        loss_list.append(loss.data)
        iteration_list.append(count)
        accuracy_list.append(accuracy)
        if count % 5 == 0:
            print('Iteration: {} Loss: {} Accuracy: {} %'.format(count,
loss.data, accuracy))

```

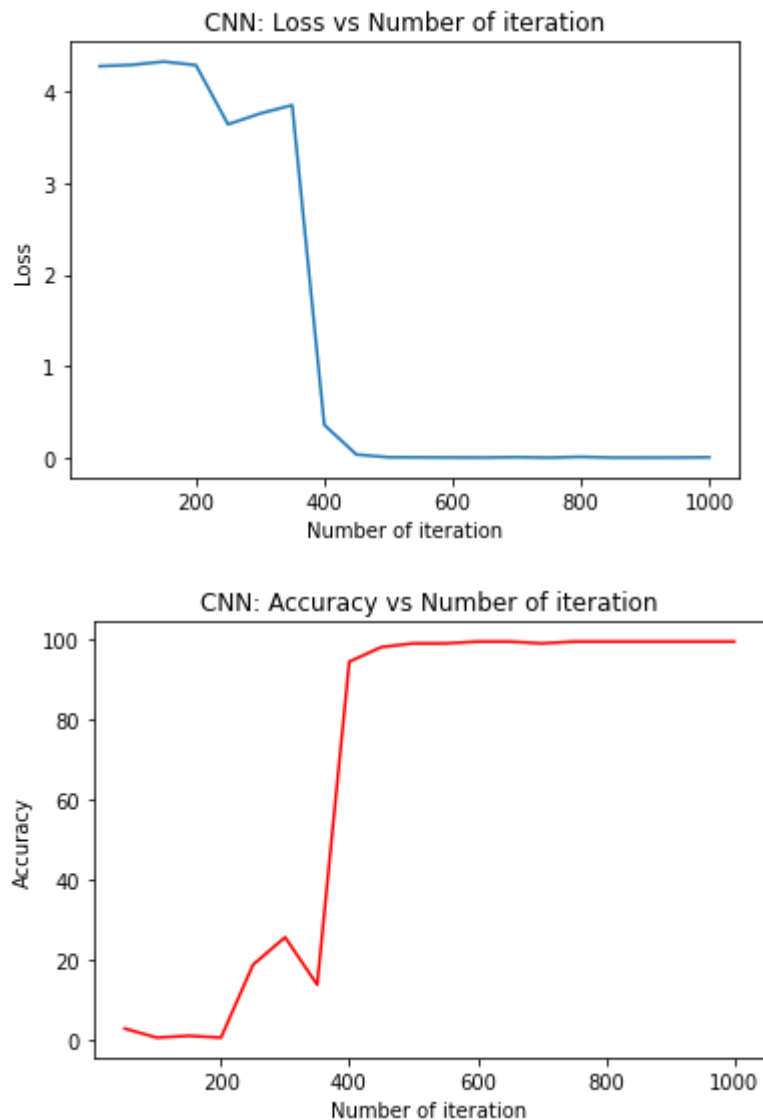
```

Iteration: 50 Loss: 4.275338649749756 Accuracy: 2.7397260665893555 %
Iteration: 100 Loss: 4.289216995239258 Accuracy: 0.456620991230011 %
Iteration: 150 Loss: 4.325140476226807 Accuracy: 0.913241982460022 %
Iteration: 200 Loss: 4.286386013031006 Accuracy: 0.456620991230011 %
Iteration: 250 Loss: 3.63972806930542 Accuracy: 18.721460342407227 %
Iteration: 300 Loss: 3.7585675716400146 Accuracy: 25.570775985717773 %
Iteration: 350 Loss: 3.8496618270874023 Accuracy: 13.698630332946777 %
Iteration: 400 Loss: 0.3587438464164734 Accuracy: 94.52054595947266 %
Iteration: 450 Loss: 0.03490452095866203 Accuracy: 98.17351531982422 %
Iteration: 500 Loss: 0.004759675357490778 Accuracy: 99.08676147460938 %
Iteration: 550 Loss: 0.003552044276148081 Accuracy: 99.08676147460938 %
Iteration: 600 Loss: 0.0022787991911172867 Accuracy: 99.54338073730469 %
Iteration: 650 Loss: 0.0014323372161015868 Accuracy: 99.54338073730469 %
Iteration: 700 Loss: 0.005599956959486008 Accuracy: 99.08676147460938 %
Iteration: 750 Loss: 0.00046150441630743444 Accuracy: 99.54338073730469 %
Iteration: 800 Loss: 0.00920148566365242 Accuracy: 99.54338073730469 %
Iteration: 850 Loss: 0.0006191439460963011 Accuracy: 99.54338073730469 %
Iteration: 900 Loss: 0.0011795498430728912 Accuracy: 99.54338073730469 %
Iteration: 950 Loss: 0.001536051044240594 Accuracy: 99.54338073730469 %
Iteration: 1000 Loss: 0.004193472675979137 Accuracy: 99.54338073730469 %

```

```
In [16]: # visualization loss
plt.plot(iteration_list,loss_list)
plt.xlabel("Number of iteration")
plt.ylabel("Loss")
plt.title("CNN: Loss vs Number of iteration")
plt.show()

# visualization accuracy
plt.plot(iteration_list,accuracy_list,color = "red")
plt.xlabel("Number of iteration")
plt.ylabel("Accuracy")
plt.title("CNN: Accuracy vs Number of iteration")
plt.show()
```



CNN with 3 convolutional layer and 1 fully connected layer

```
In [17]: # Create CNN Model
class CNNModel(nn.Module):
    def __init__(self):
        super(CNNModel, self).__init__()

        self.cnn1 = nn.Conv2d(in_channels=3,out_channels=16,kernel_size=5,stri
de=1,padding=0)
        self.relu1 = nn.ReLU()

        self.maxpool1 = nn.MaxPool2d(kernel_size=2)

        self.cnn2 = nn.Conv2d(in_channels=16,out_channels=32,kernel_size=5,str
ide=1,padding=0)
        self.relu2 = nn.ReLU()

        self.maxpool2 = nn.MaxPool2d(kernel_size=2)

        self.cnn3 = nn.Conv2d(in_channels=32,out_channels=64,kernel_size=5,str
ide=1,padding=0)
        self.relu3 = nn.ReLU()

        self.maxpool3 = nn.MaxPool2d(kernel_size=2)

        self.fc1 = nn.Linear(5184,len(target_labels));

    def forward(self,x):

        out=self.cnn1(x)
        out=self.relu1(out)
        out=self.maxpool1(out)

        out=self.cnn2(out)
        out=self.relu2(out)
        out=self.maxpool2(out)

        out=self.cnn3(out)
        out=self.relu3(out)
        out=self.maxpool3(out)

        out = out.view(out.size(0), -1)
        out = self.fc1(out)

        return out
```

```
In [18]: # Initialize Parameters and fit the model
model = CNNModel()
error = nn.CrossEntropyLoss()
learning_rate = 0.02
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```

```

In [19]: # CNN model training
count = 0
loss_list = []
iteration_list = []
accuracy_list = []
for epoch in range(num_epochs):
    for i, (images, labels) in enumerate(train_loader):
        train = Variable(images.view(-1,3,100,100))
        labels = Variable(labels)
        optimizer.zero_grad()
        outputs = model(train)
        loss = error(outputs, labels)
        loss.backward()
        optimizer.step()
        count += 1

    if count % 50 == 0:
        correct = 0
        total = 0
        for images, labels in test_loader:
            test = Variable(images.view(-1,3,100,100))
            outputs = model(test)
            predicted = torch.max(outputs.data, 1)[1]
            total += len(labels)
            correct += (predicted == labels).sum()
        accuracy = 100 * correct / float(total)
        loss_list.append(loss.data)
        iteration_list.append(count)
        accuracy_list.append(accuracy)
        if count % 5 == 0:
            print('Iteration: {} Loss: {} Accuracy: {} %'.format(count,
loss.data, accuracy))

```

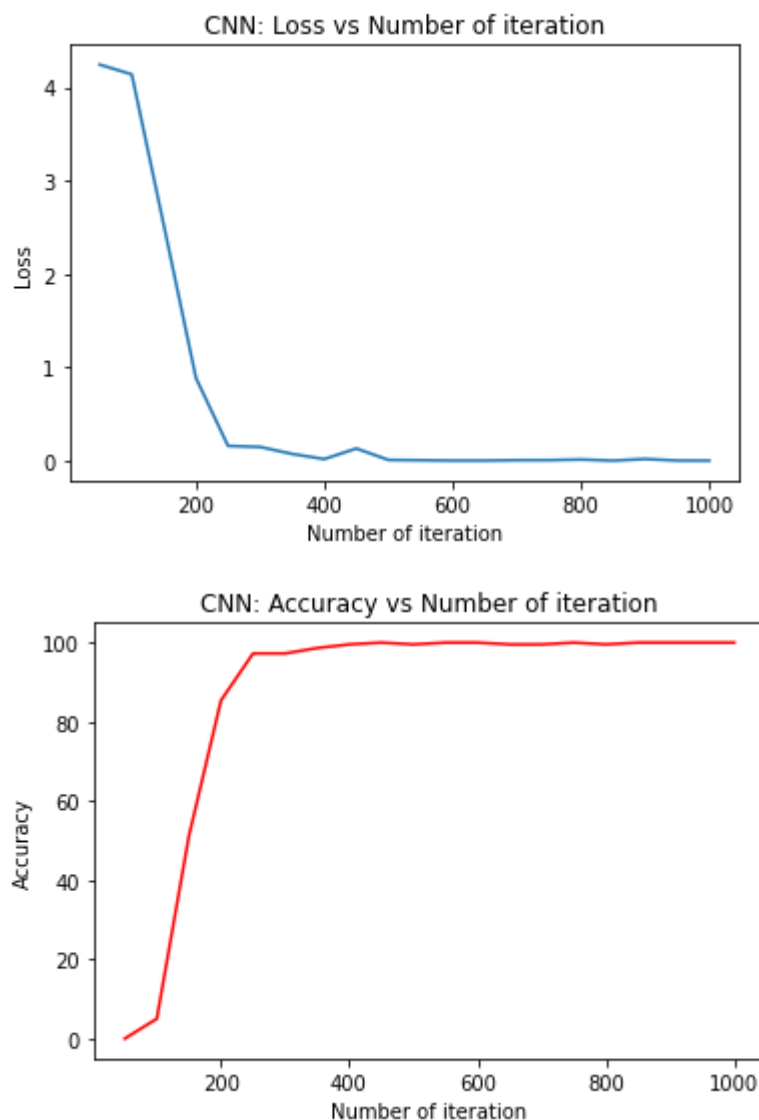
```

Iteration: 50 Loss: 4.2426676750183105 Accuracy: 0.0 %
Iteration: 100 Loss: 4.137964248657227 Accuracy: 5.022830963134766 %
Iteration: 150 Loss: 2.5207266807556152 Accuracy: 51.14155197143555 %
Iteration: 200 Loss: 0.8878799676895142 Accuracy: 85.38813018798828 %
Iteration: 250 Loss: 0.16059371829032898 Accuracy: 97.2602767944336 %
Iteration: 300 Loss: 0.14908604323863983 Accuracy: 97.2602767944336 %
Iteration: 350 Loss: 0.0743330642580986 Accuracy: 98.63013458251953 %
Iteration: 400 Loss: 0.019346822053194046 Accuracy: 99.54338073730469 %
Iteration: 450 Loss: 0.1324983835220337 Accuracy: 100.0 %
Iteration: 500 Loss: 0.010032325983047485 Accuracy: 99.54338073730469 %
Iteration: 550 Loss: 0.006634755525738001 Accuracy: 100.0 %
Iteration: 600 Loss: 0.0030088969506323338 Accuracy: 100.0 %
Iteration: 650 Loss: 0.0032461590599268675 Accuracy: 99.54338073730469 %
Iteration: 700 Loss: 0.007285582832992077 Accuracy: 99.54338073730469 %
Iteration: 750 Loss: 0.0077935331501066685 Accuracy: 100.0 %
Iteration: 850 Loss: 0.0024780724197626114 Accuracy: 100.0 %
Iteration: 900 Loss: 0.01989869400858879 Accuracy: 100.0 %
Iteration: 950 Loss: 0.0033794406335800886 Accuracy: 100.0 %
Iteration: 1000 Loss: 0.0025697543751448393 Accuracy: 100.0 %

```

```
In [20]: # visualization loss
plt.plot(iteration_list,loss_list)
plt.xlabel("Number of iteration")
plt.ylabel("Loss")
plt.title("CNN: Loss vs Number of iteration")
plt.show()

# visualization accuracy
plt.plot(iteration_list,accuracy_list,color = "red")
plt.xlabel("Number of iteration")
plt.ylabel("Accuracy")
plt.title("CNN: Accuracy vs Number of iteration")
plt.show()
```



As we can see from above two model , if first I take 2 convolutional layer and 1 full connected layer with learning rate 0.1, next I take 3 convolutional layer and 1 fully connected layer with learning rate 0.02, keeping all parameters same I get good performance but If I take other learning rate with different numbers of convolutinal layer and fully connected layer then I get bad performance or underfitting model.